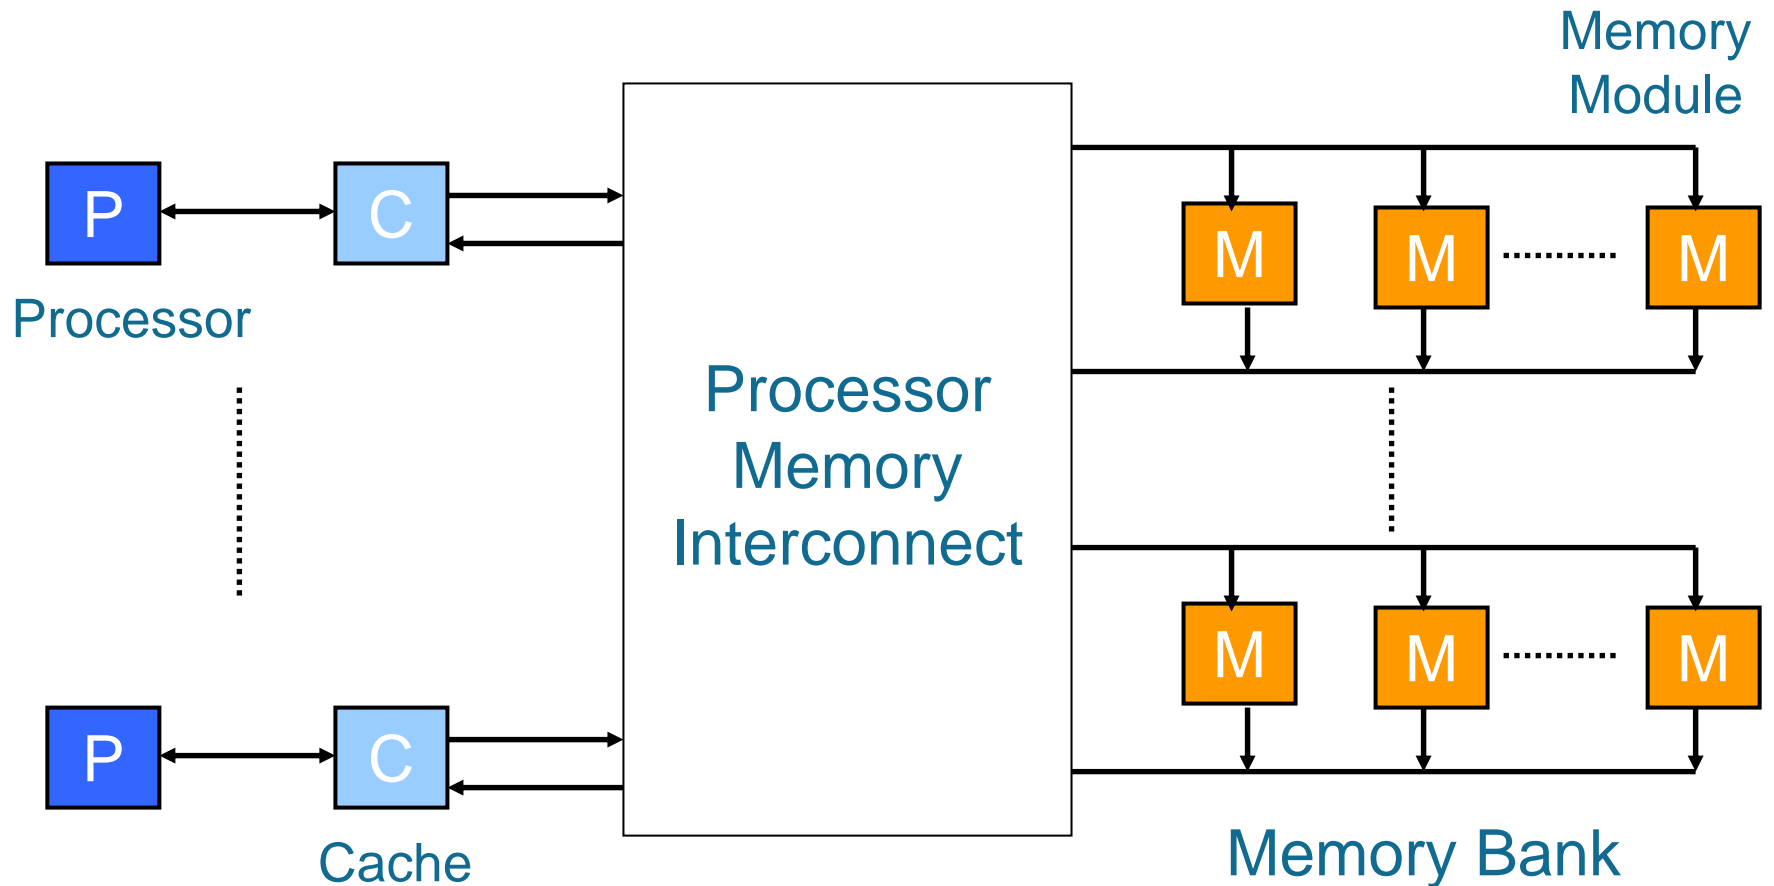


ARM Architecture


Memory Subsystem

Ashish Kuvelkar
Hardware Technology Development Group
C-DAC, Pune

Canonical Computer System



Components of a Computer



Memory System Performance

- Access time
 - Time elapsed between time request is made and data reaches processor
 - Dependent on bus delay, chip access delay
- Memory Bandwidth
 - Amount of information transferred over a period of time. Depends upon
 - Way physical memory is organised
 - Number of banks
 - Accessing modes within module

Memory Requirements

- Programmers want unlimited amount of fast memory
- Hardware aids programmer by creating an illusion of unlimited fast memory
- Realities
 - Fast memory is costly
 - Programs tend to follow principle of locality
 - Temporal: same location is repeatedly accessed
 - Spatial: Close by locations are accessed soon

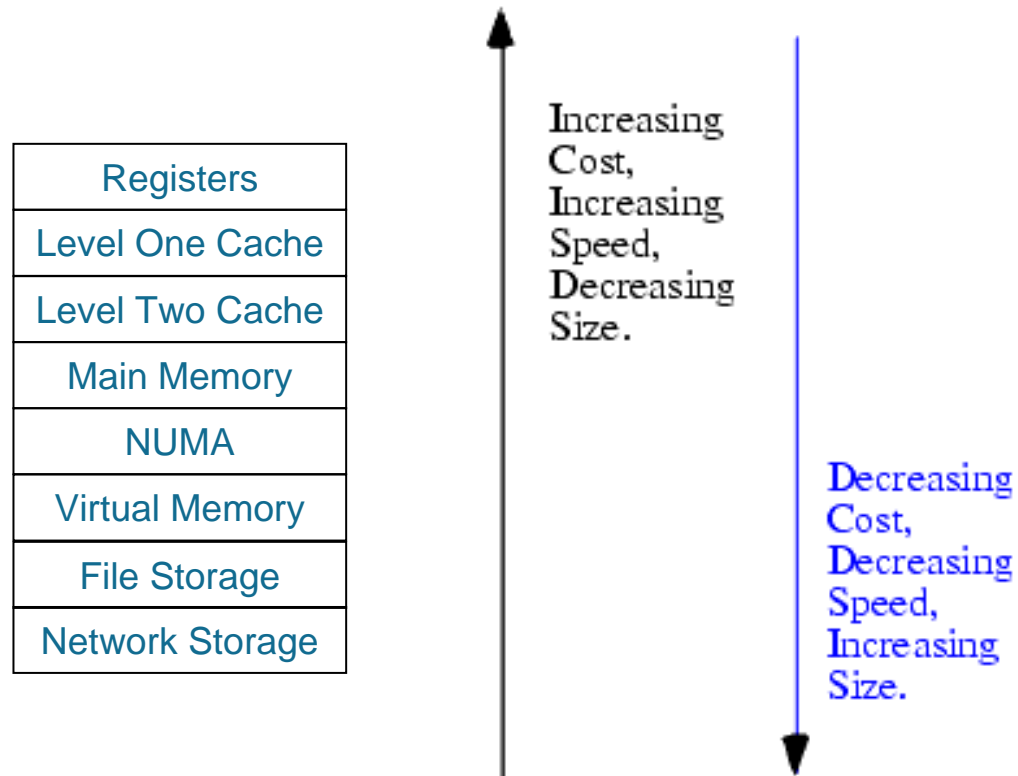
Processor and memory speeds

- When RISC processors were introduced, standard memory parts were faster than contemporary microprocessors
- Over a period
 - Microprocessors have become faster
 - Memories have become little faster but mostly have higher capacity
- The memories are not able to keep up with the speed achieved by microprocessors

Memory Hierarchy


- Takes advantage of principle of locality
- Memory subsystem consists of multiple levels with different speeds and sizes.
- Small fast memory is put close to processor and slower away
- The goal
 - Provide maximum memory which is cheapest
 - Provide access at the speed offered by the fastest
- Data is copied only between two adjacent levels
- If data is found in level closest to CPU, it's a hit
- Hit rate is a measure of the performance

Memory Hierarchy



Cache

Cache: as safe place for hiding or storing things
-Webster's New world dictionary

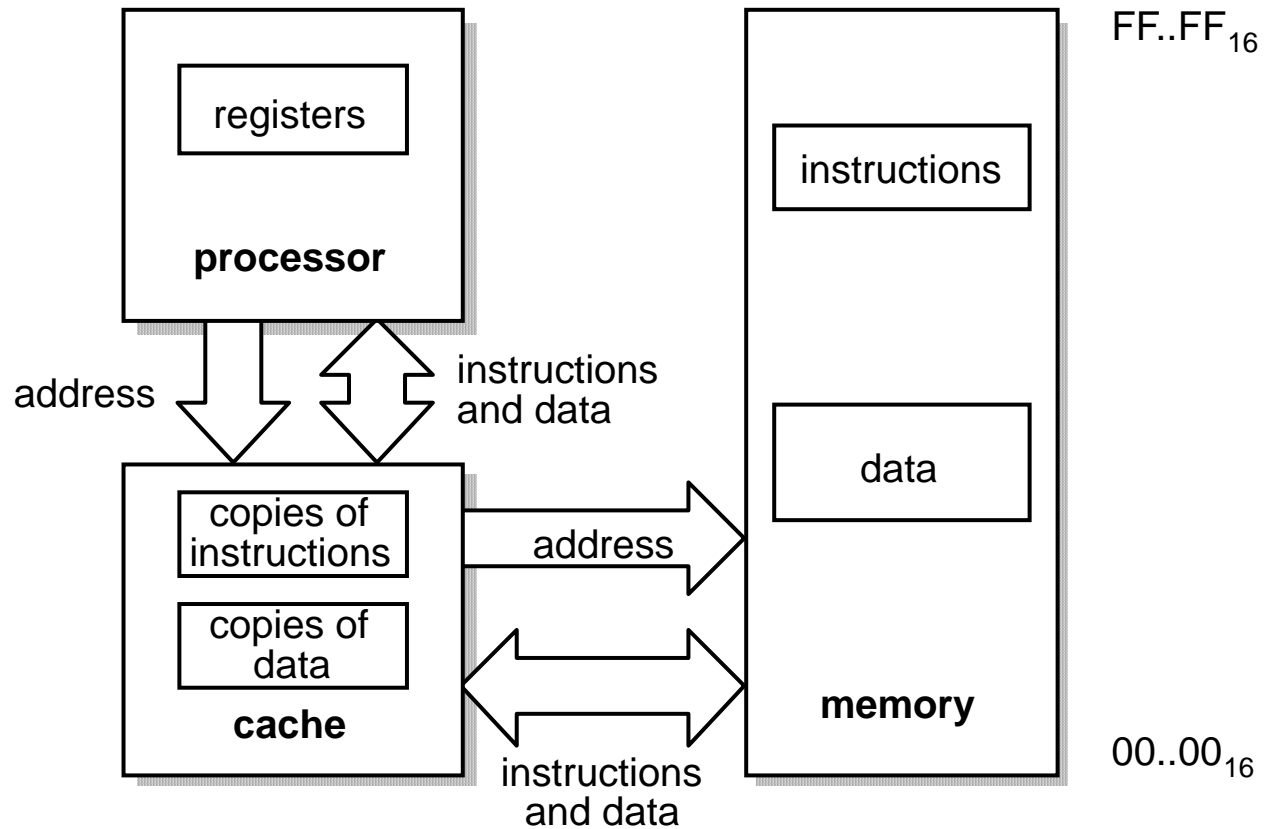


- Cache contains a copy of portions of main memory
- When CPU attempts to read a word, cache is checked first. If found, it is presented, else a line of words is brought from main memory to the cache

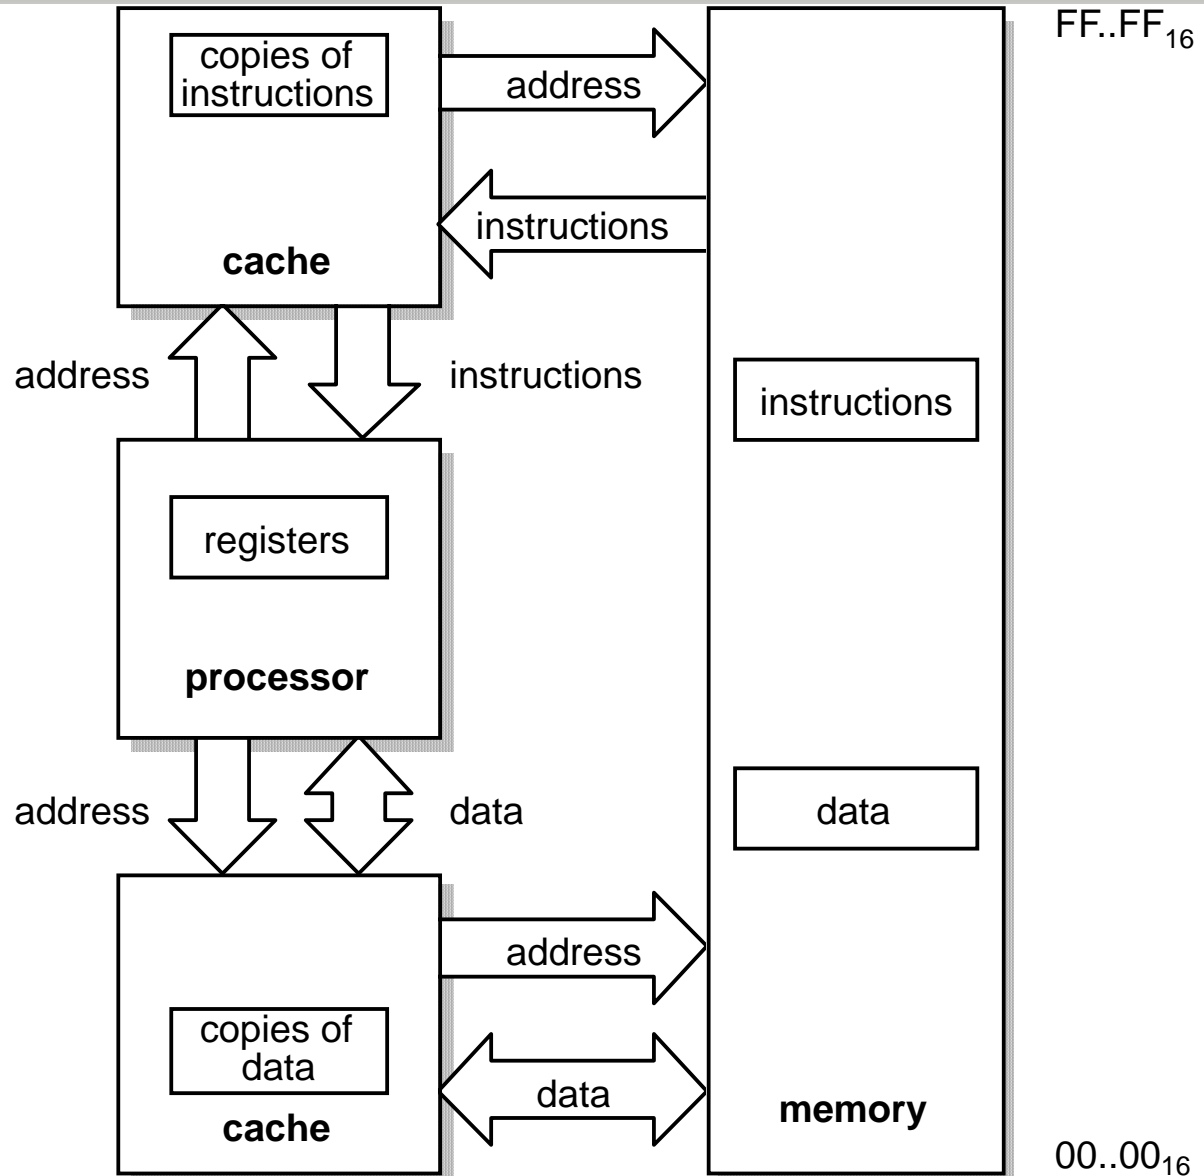
Cache Organisation

- At the highest level the processor can have two ways of organising Cache
 - Unified Cache
 - Automatically adjust the proportion of cache memory for instructions, according to current programming requirements
 - Separate instruction and data cache
 - Allows load and store instructions to execute in single clock cycle

Unified Instruction and data cache



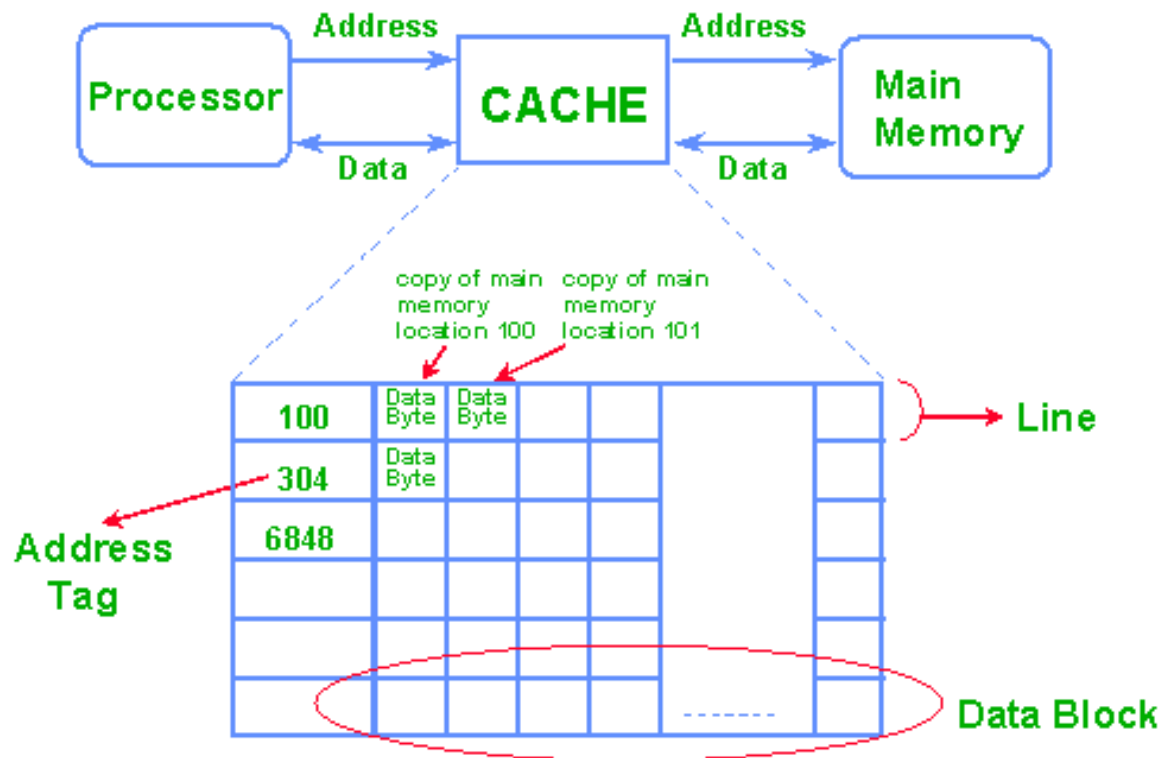
Separate Instruction and data caches



Basic Architecture of Cache

- A cache consists of 3 main parts
 - Directory store
 - The cache must know where the information in cache originated from
 - Each directory entry is called as cache-tag
 - Data section
 - Holds the data copied from main memory
 - Status information
 - Valid bit indicates live data
 - Dirty bit indicates non-coherency

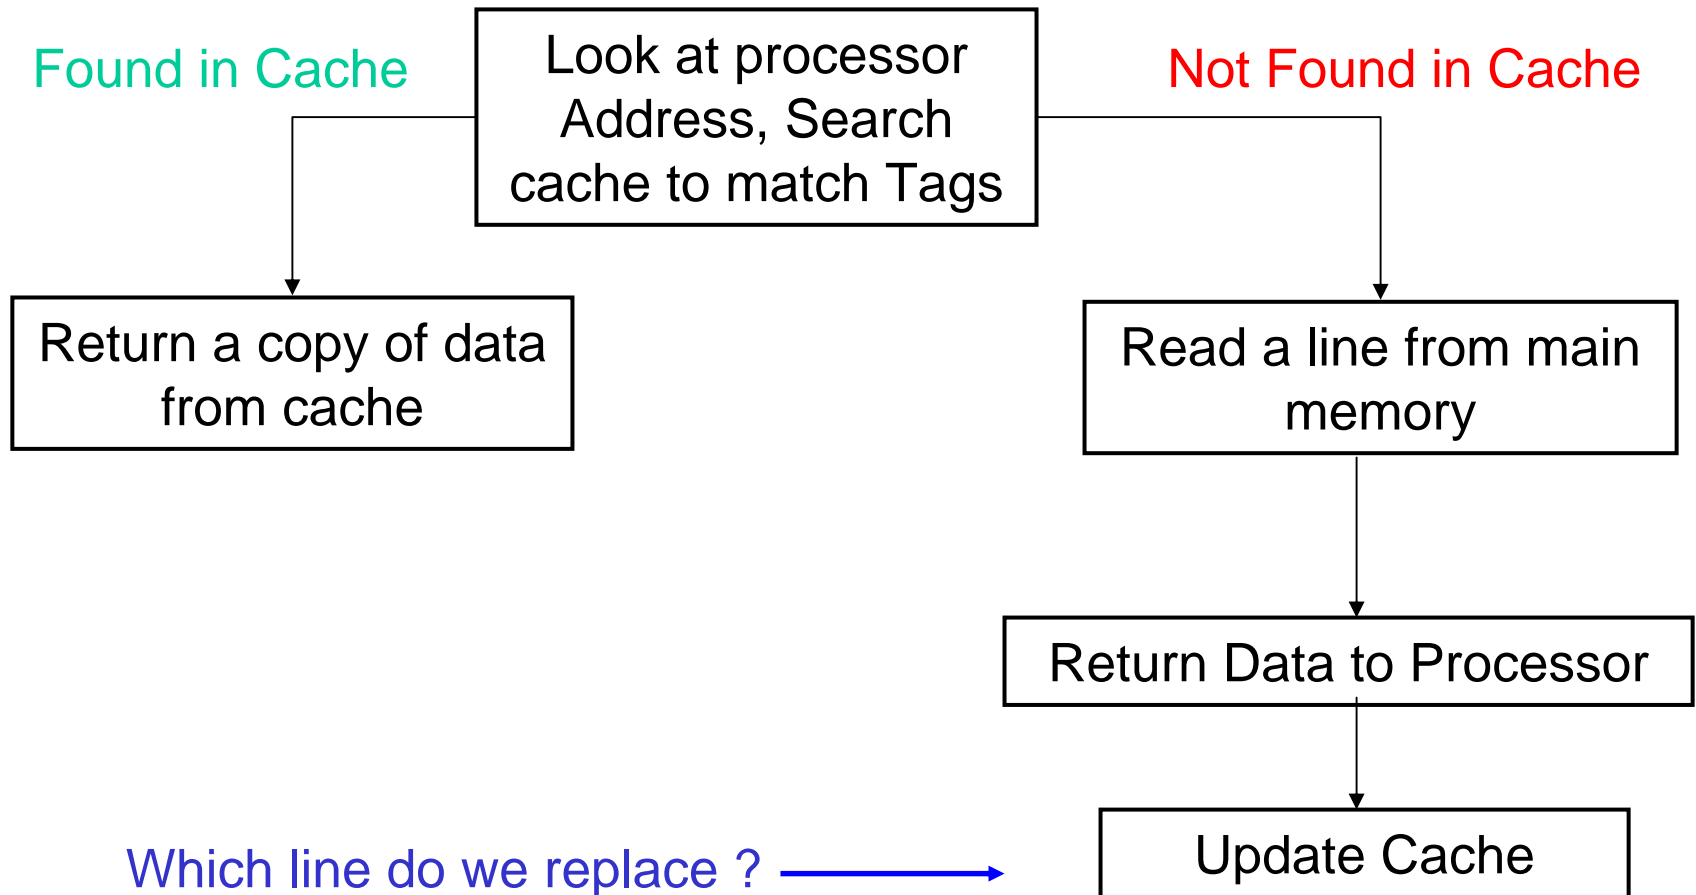
Inside a Cache



Cache Controller

- It's the hardware that copies code or data from main memory to cache memory automatically.
- The operation is transparent to CPU and the software
- It intercepts reads and writes before passing them on to the memory controller.
- The relationship between main memory and cache is called as mapping

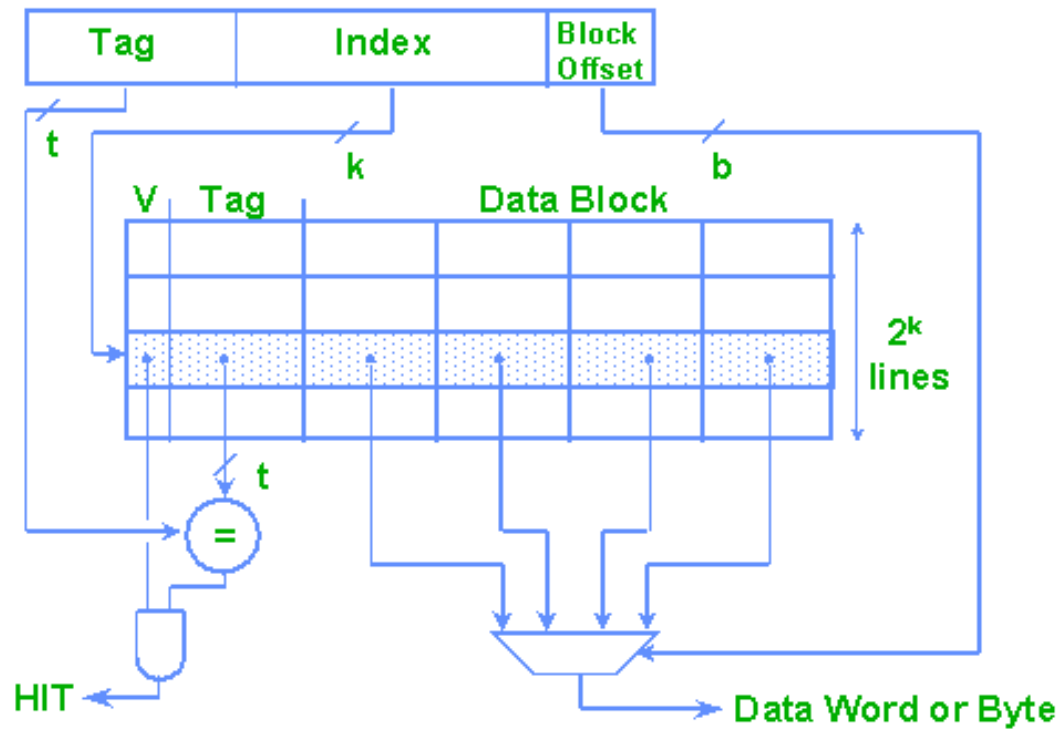
Cache Read Algorithm



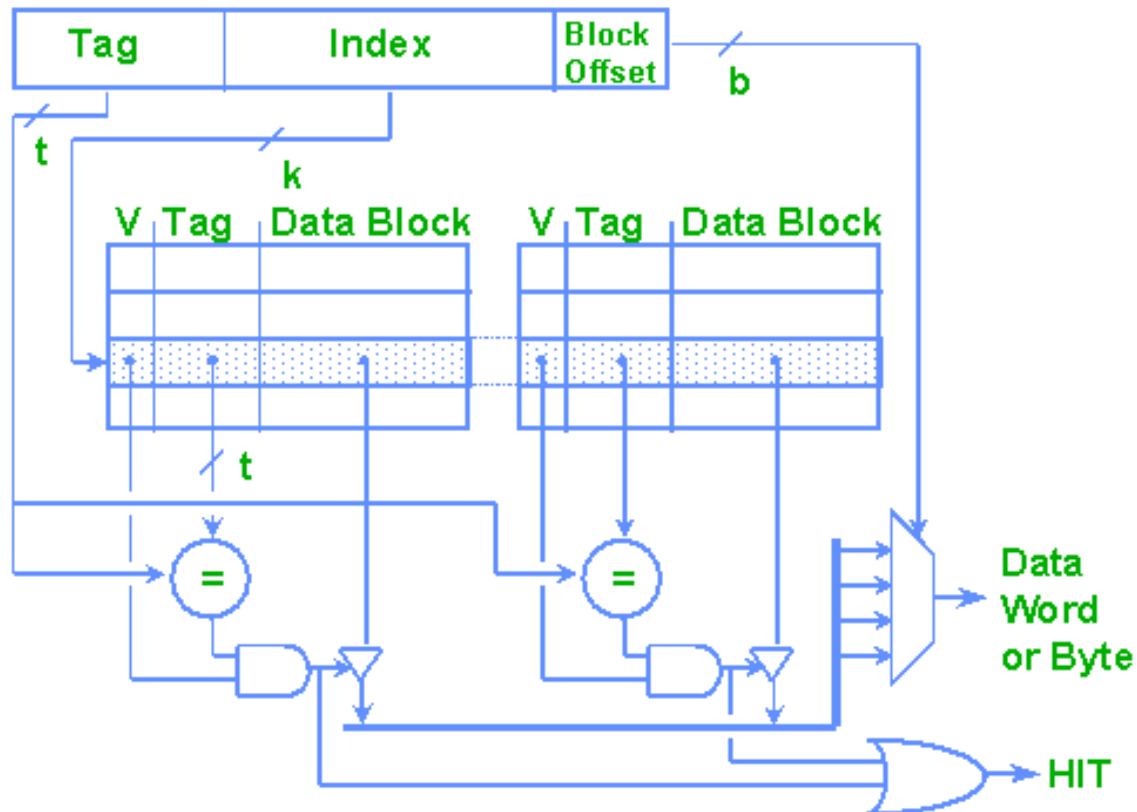
Cache Mapping

- Mapping between main memory locations and cache
 - Direct mapped
 - N-way set-associative ($n=2,4$)
 - Fully set-associative

Direct mapped



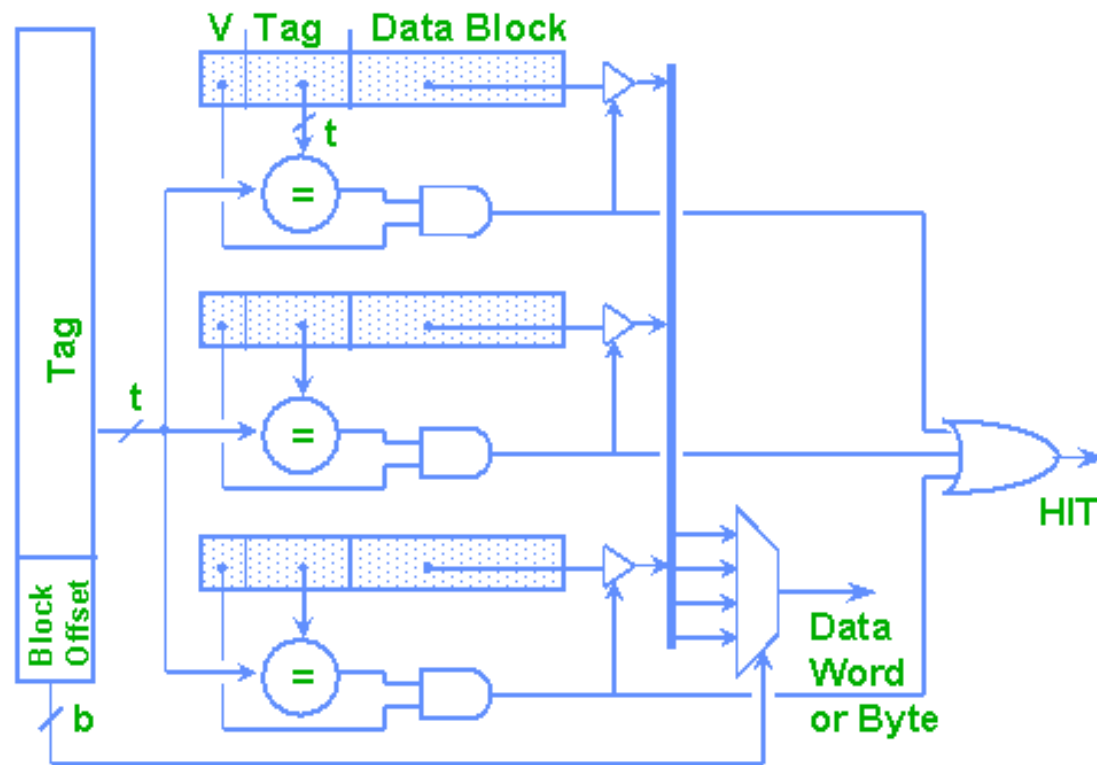
2-way Set Associative Cache



2-way Set Associative Cache

- Improved direct mapped cache
- Each memory line has two destinations in cache
- Extra logic is required for comparing tags
- Replacement policy needs consideration

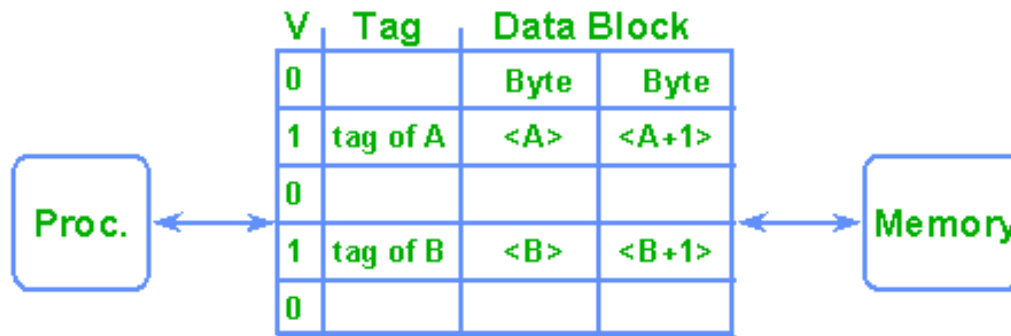
Fully Associative Cache



Fully Associative Cache

- This scheme represents other extreme
 - Any line can be placed in any location in cache
- Gives highest hit ratio
- Search is slowest
- Content addressable Memories (CAM) are used for storing Tags

Valid bits



- Valid bit must be **1** for cache line to **HIT**.
- At power-up or reset, we set all valid bits to **0**.
- Set valid bit to **1** when cache line is first replaced.
- Flush cache by setting all valid bits to **0**, under external program control.

Cache Efficiency

- Two terms are use to characterise cache efficiency of a program
 - Cache hit rate
 - Cache miss rate

$$\text{Hit rate} = \frac{\text{cache hits}}{\text{memory requests}} \times 100$$

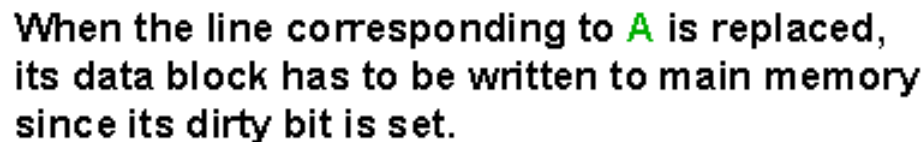
- Rates can measure read, writes or both
- Other performance measurement terms are hit time and miss penalty

Cache Policy

- Three policies determine cache operation
 - Write policy
 - Determines where data is stored during processor write operation
 - Replacement policy
 - Determines the line that will be used for next line fill during a cache miss
 - Allocation policy
 - Determines when the cache controller allocates a line

Write Policy

- Two alternatives when processor issues a write
 - Update both cache and main memory: **write-through**
 - Update only the cache: **write-back**
- Write-through
 - Both memories remain coherent all the time
 - Writes are slower
- Write-back
 - Valid cache lines and memory may not be coherent
 - Dirty bits are used to indicate non-coherency
 - Dirty cache line are updated before eviction



Replacement Policy

- After a read miss, which cache line should be replaced with the data line read from main memory
 - Direct mapped cache: unique line
 - Associative cache
 - Least Recently Used(LRU): Replace line in set corresponding to address which has not been read or written for the longest time
 - First-in First Out(FIFO): Select line which has been in the set for the longest time
 - Random: Select line in the set randomly

Allocation Policy

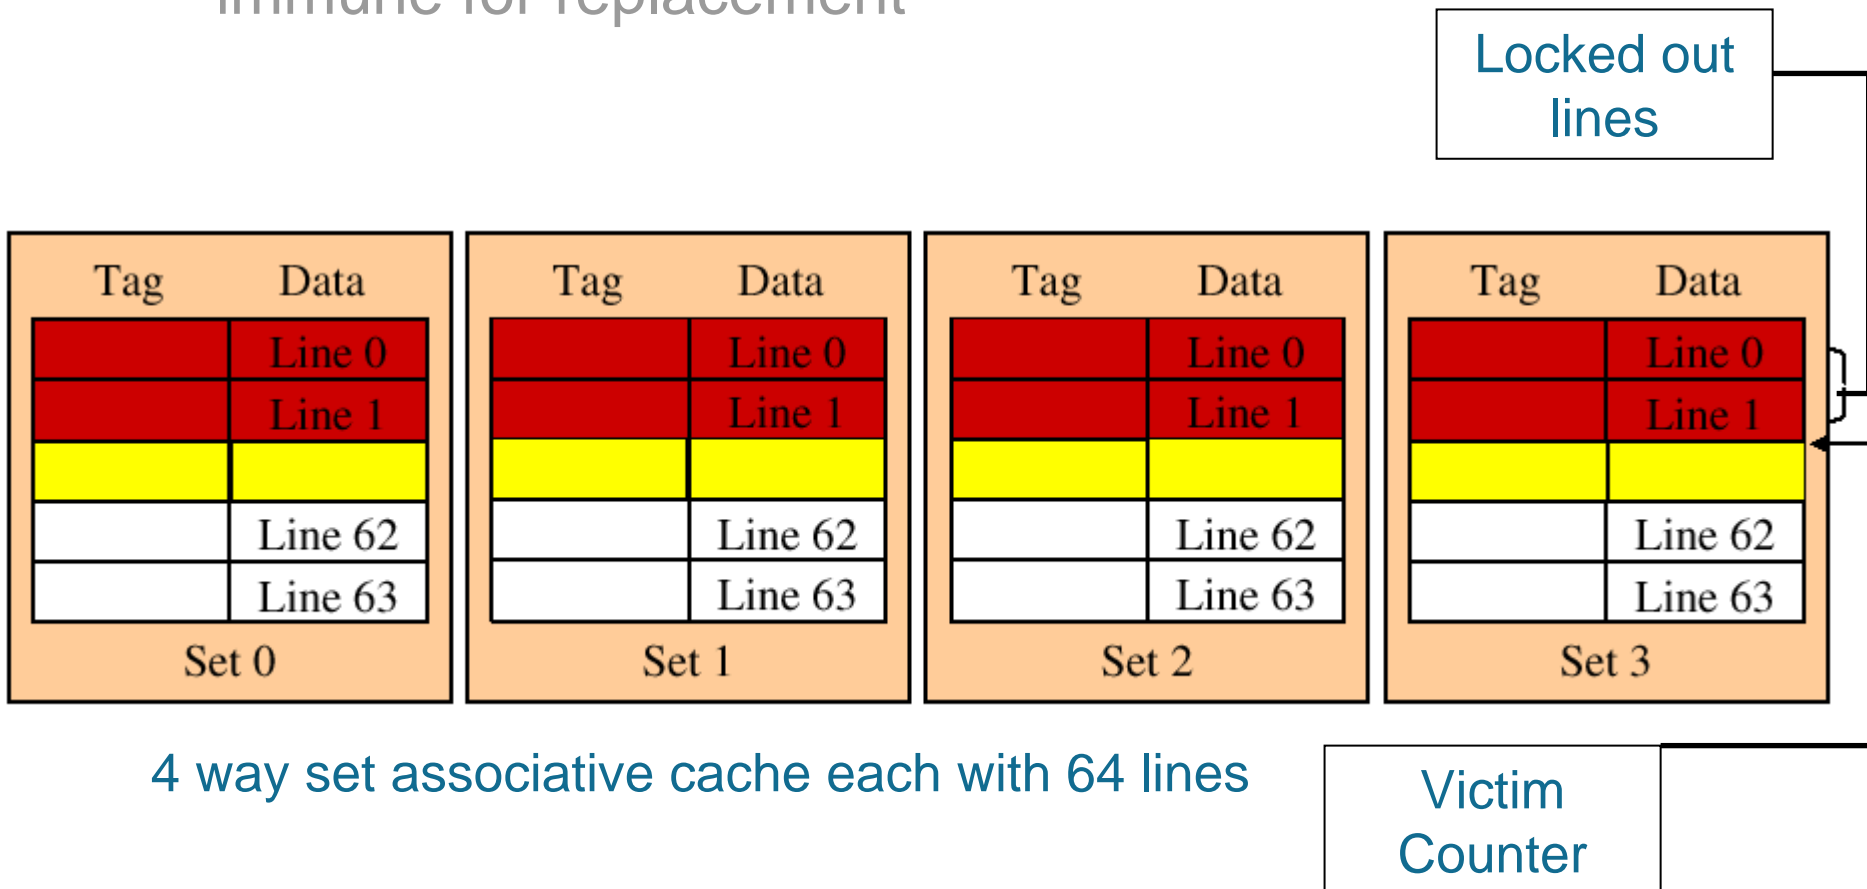
- Two strategies to allocate cache line on a cache miss
 - Read-allocate
 - Allocation is done during read from main memory
 - A write does not update cache unless a line was allocated on previous read
 - Read-write-allocate
 - Allocation is done during either a read or write
 - On write, if cache line is not valid, it does a cache fill before updating the cache.
 - Write policy determines whether main memory should be updated

Cache Lockdown

- A feature that enables a program to load time-critical code and data into cache and exempt it from eviction
- It avoids the problem of unpredictable execution times that result from line replacement
- Cache used for lockdown results in reduction of cache available for other parts of main memory
- Candidates for locking in cache
 - Vector interrupt table
 - Interrupt service routines
 - Critical code used frequently
 - Frequently used global variables

Cache lockdown

- Lines below chosen by victim counter are immune for replacement



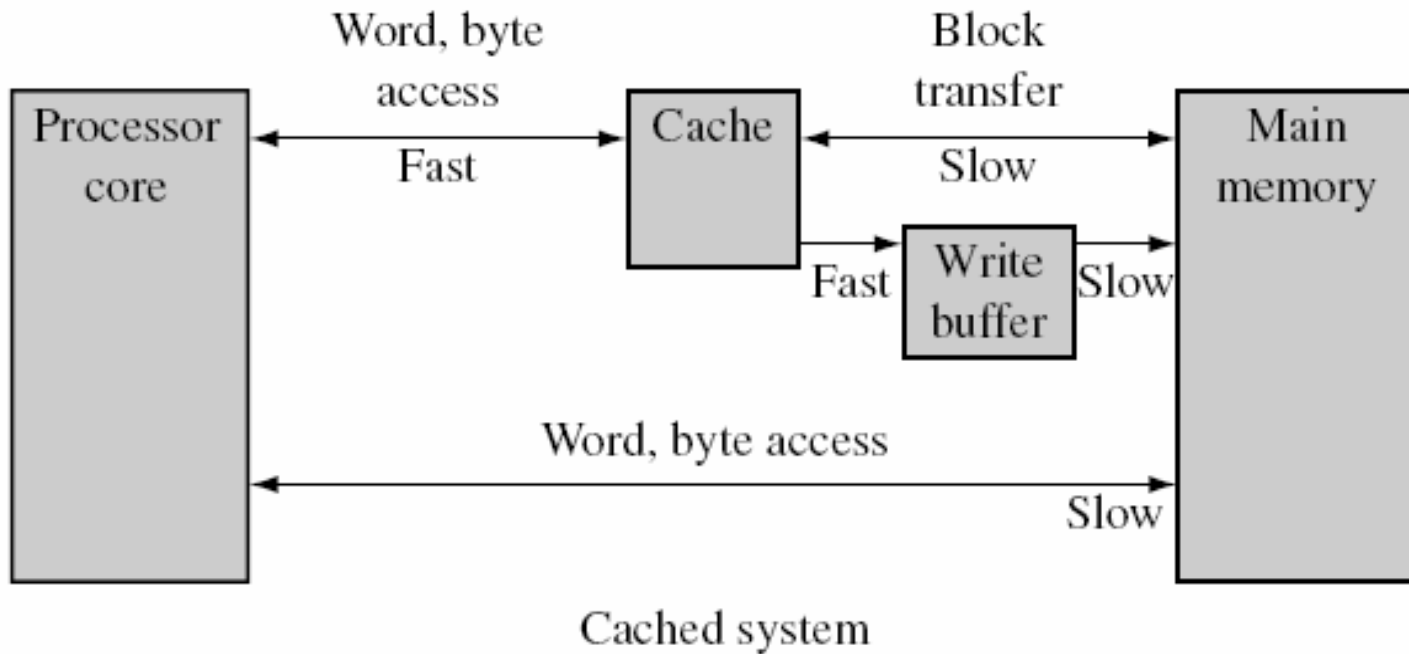
Write Buffers

- Ensures that memory is decoupled from core
 - Data placed in buffer at core speed
 - In parallel data written to memory at bus speed
- Access are regulated to occur in correct order
 - Dedicated write buffer
 - Write buffer is always drained first for non cached reads, cache line fills and non-buffered writes


Tightly coupled memories

- Fast memory local to processor
 - High speed performance without access to system bus
 - Real time performance can be predicted
- Alternative to cache
 - Smaller die size and low power
- Appears at fixed locations in memory area
 - Some cores have both TCM and cache, some TCM
 - Different TCMs for instruction and data

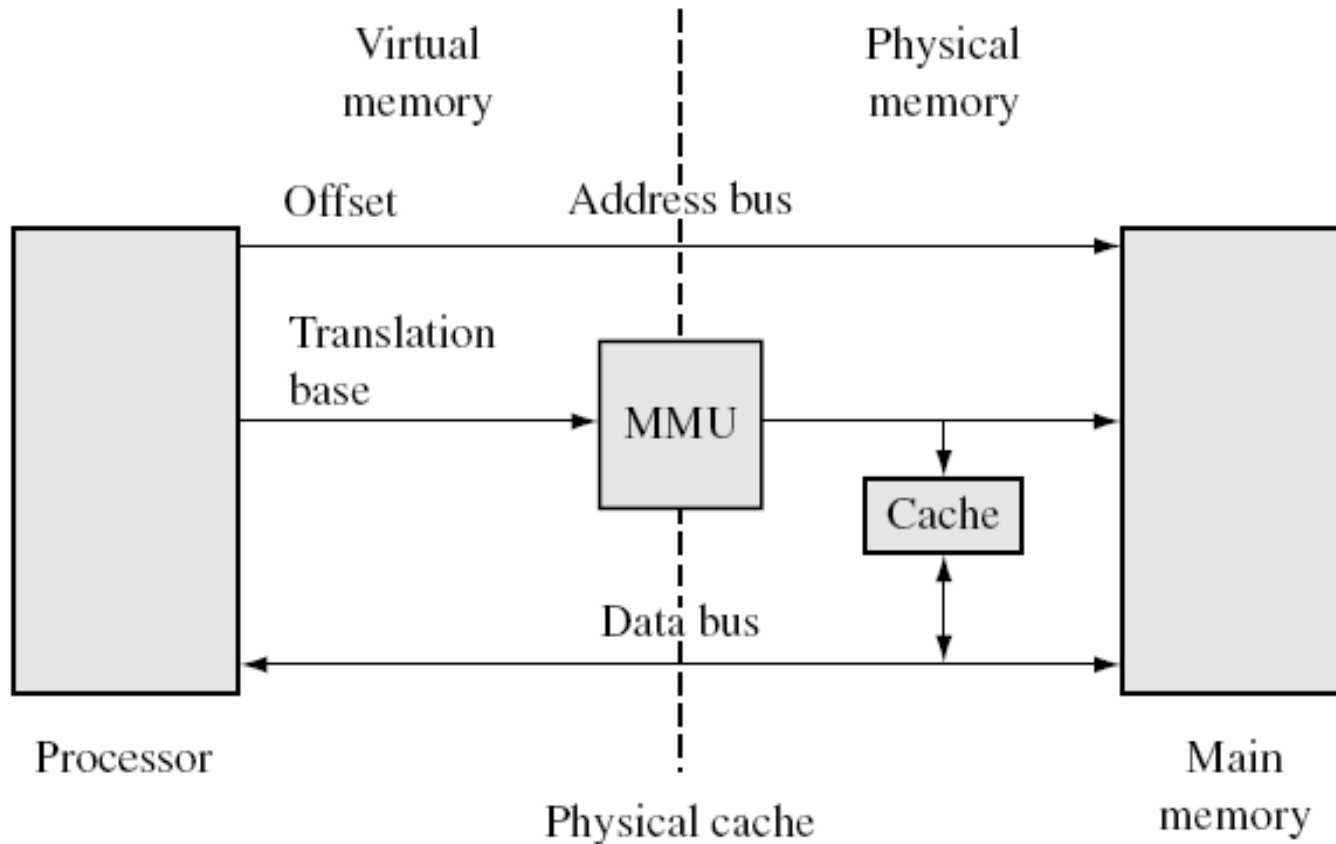
Write Buffers



Logical Cache



Physical Cache




Main Memory

DRAMs

- Over the years evolution of system memory has taken place
 - Asynchronous DRAM technologies,
 - Fast Page Mode (FPM) memory and
 - Extended Data Out (EDO) memory,
 - synchronous DRAM (SDRAM)
 - Registered SDRAM
 - Double Data Rate (DDR) SDRAM
 - DDR – II SDRAM
 - Rambus DRAM (RDRAM)
- Yet, system memory bandwidth has not kept pace with improvements in processor performance

Basic DRAM operation



Asynchronous DRAMs

- Address is provided in two parts qualified by
 - Row Address Strobe (RAS)
 - Column Address Strobe (CAS)
- Data is available after CAS is asserted
- Memory controller asserts signals to suit DRAM timing specifications
- Latency is longer due to transfer of data between asynchronous DRAMs and synchronous system bus

FP and EDO DRAMs

- All locations having same row address is a page
- FP DRAMs allows any location in current page to be access by strobing in only column address with a CAS.
- This reduces latency for subsequent accesses
- EDO DRAMs maintain the output for longer period even when CAS is de-asserted.
- This improves overall system efficiency

Synchronous DRAMs

- Memory bus clock is used to synchronize the input and output signals on the memory chip
 - Simplifies the memory controller
 - Reduces the latency from CPU to memory
- Memory is divide into two to four banks for simultaneous access to more data
- The bandwidth capacity of the memory bus increases with its width (in bits) and its frequency (in MHz).
- For a system transferring 8 bytes (64 bits) at a time and running at 100 MHz, the bandwidth is 800 MB/s,

Memory Management and Protection

Memory Addresses

- The CPU supports a variety of addressing modes
- The address generated by address unit of CPU refers to a logical space of contiguous locations
- The memory presents itself as a contiguous space of physical locations
- A relationship needs to be established between the logical address space of CPU and the physical address space presented by memory.
- In CPUs like 8085, the mapping between the two is one is to one.

Constraints for memory management

- User programs need logical space that is much bigger than the available physical memory
- Assignment and sub-division of physical address space to a number of different user programs
- The protection of the physical address area assigned to a user from possible interference by other user programs

Solutions

- To allow large logical address space, virtual memory used.
 - Address space is mapped onto secondary storage like disks
 - Locality principle helps in keeping frequently used program segments into main memory
- To allow assignment of physical memory to programs dynamic relocation is used
 - Logical addresses are assigned physical address at the time of loading

Solutions

- To prevent unauthorized programs from accessing memory areas
 - Each physical memory area is provided with access attributes
 - Read, write, execute etc

Mapping

- It sets the relationship between logical and physical addresses
- In the simplest of schemes
 - The complete program is treated as single area
 - It is allotted a single contiguous area in physical memory
 - Mapping parameters are base address and size
 - $\text{Physical address} = \text{base} + \text{logical address}$
 - This method has problems related to compacting and efficient use of physical memory

Dynamic mapping

- Program is treated as number of separate address areas inside which contiguity is observed
- In main memory only those parts of program are maintained which are 'active'
- This scheme needs description of where each segment of logical address is placed in physical memory
- This is achieved by using two-component address
 - The address area
 - The offset

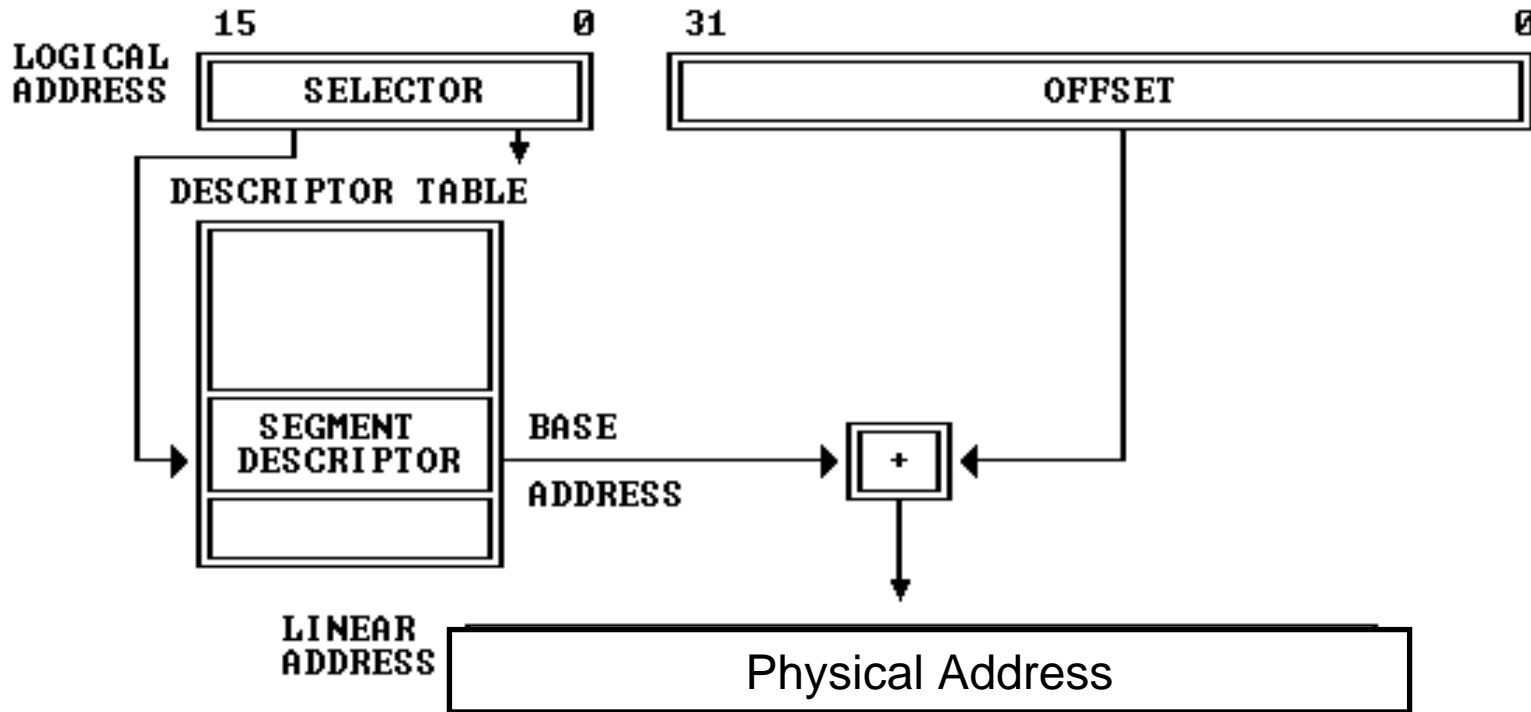
Mapping Methods

- Segmentation
 - The logical address space is broken into units called segments
 - Each segment has base address, Size and attributes
- Paging
 - Memory is divided into pages of constant length
 - Pages are moved between logical and physical address space
- Tables are maintained to achieve the mapping

Segmentation

- To perform translation, the processor uses the following data structures:
 - Segment Descriptors
 - Descriptor tables
 - Selectors
 - Segment Registers

Segment Translation



Segment Descriptors

- The segment descriptor provides the processor with the data it needs to map a logical address into a linear address.
- Descriptors are created by compilers, linkers, loaders, or the operating system
- Descriptors are not created by applications programmers.

Segment Descriptor

Each Segment is represented by a 8 byte segment descriptor

Attributes	Base (16 – 23)	Base (16-31)
Base (0-15)		Limit (0-15)

32 bit Base field : contains linear address of the first byte of the segment

24 bit limit field : denotes segment length


Attributes: Protection, privilege levels etc

Paging Unit

Paging

- Both logical and physical address space is divided into fixed sized components called pages
- Relationship between logical and physical pages is stored in page tables.
 - Single table
 - Requires very large table
 - Two or more levels of page tables

Paging memory management



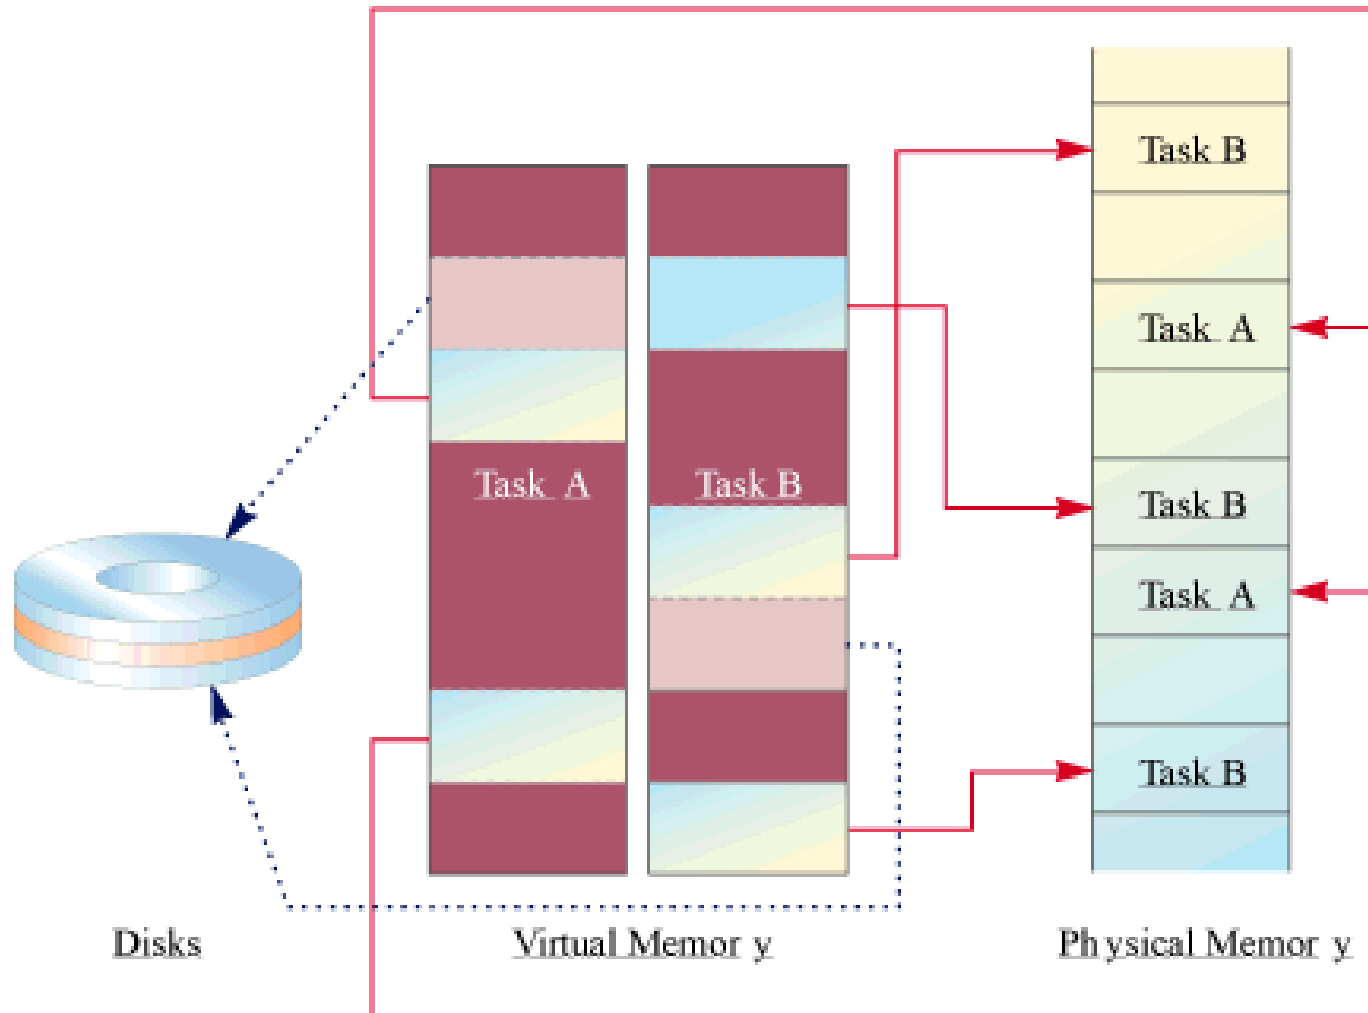
Virtual Memory

- Two motivations
 - To allow efficient sharing of memory among multiple programs
 - A collection of programs are running at a time
 - Total memory requirement of all would be more
 - Fraction of this is being used at a time
 - To remove the programming burden of small, limited amount of main memory
 - Programmers divided large programs into pieces and identified mutually exclusive ones
 - These overlays were loaded and unloaded during execution

Virtual Memory

- Functions of Virtual Memory subsystem
 - Translation
 - From a program's own address space to physical memory address
 - Protection
 - Same physical memory is shared between various programs and the OS


Features




Address Translation

Virtual Addresses

Physical Addresses

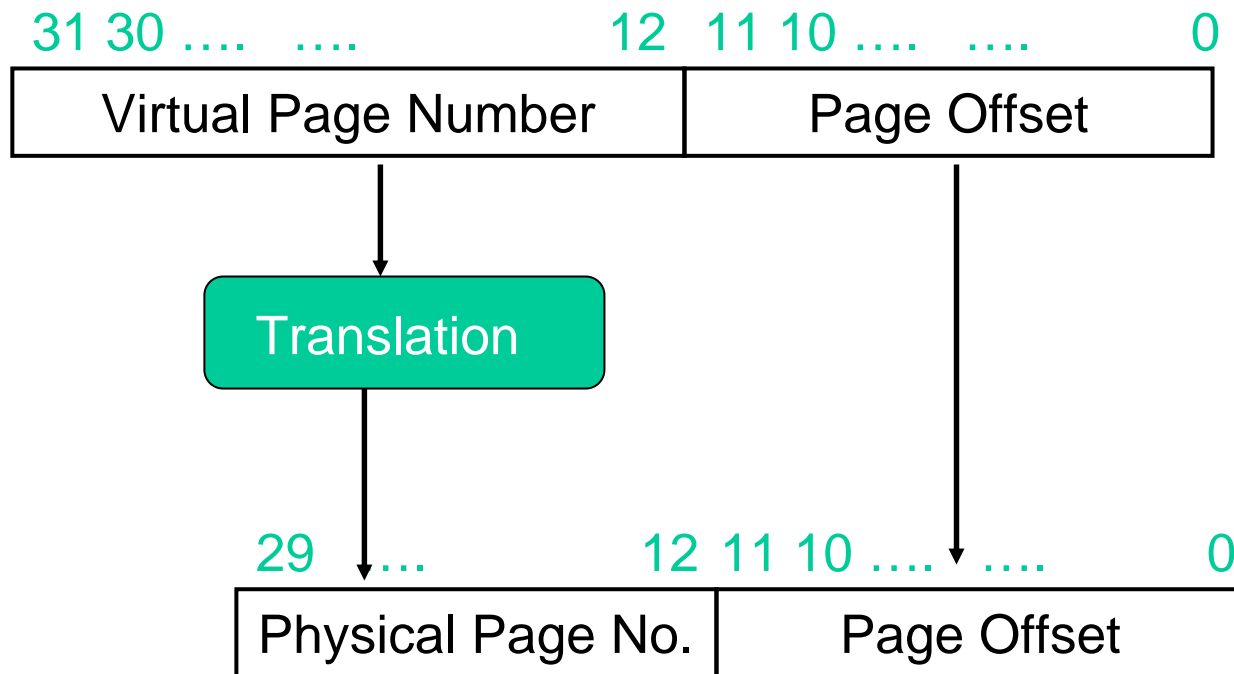


Disk Addresses



Mapping

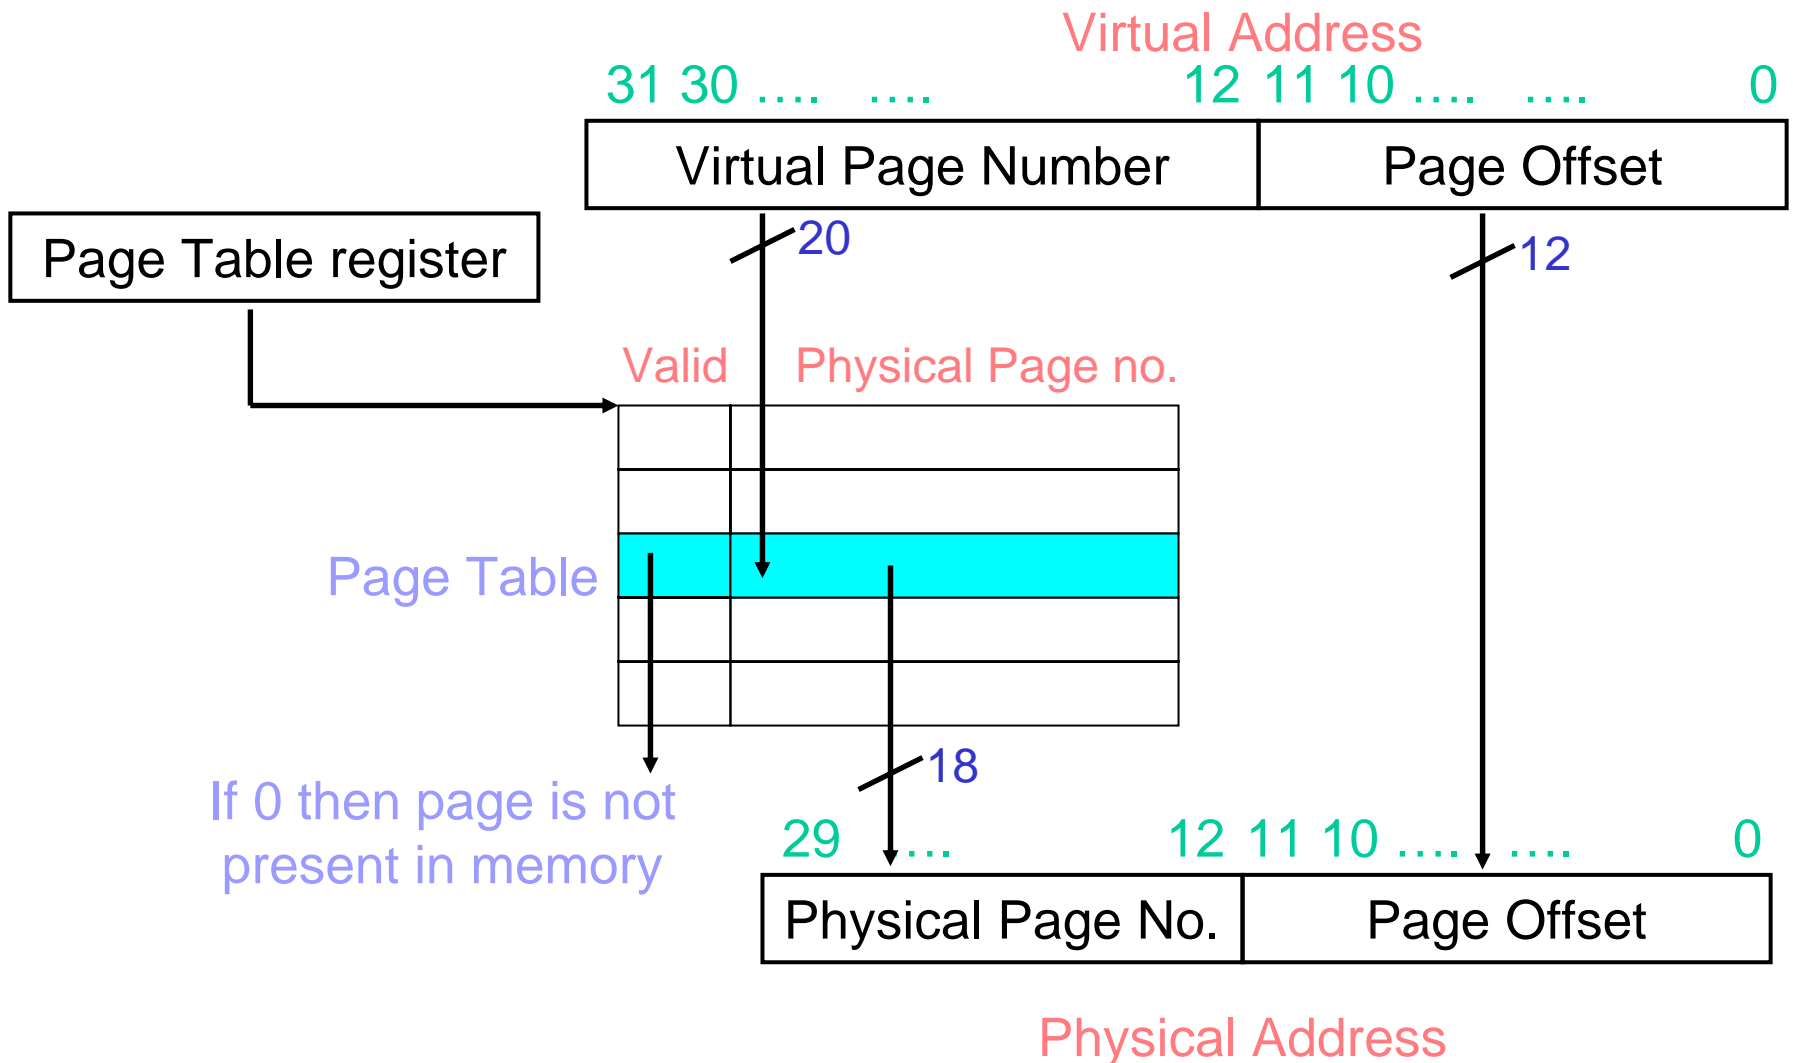
Virtual Address



Physical Address

- Page size is 4Kb
- Physical pages allowed is 2^{18}
- Main memory is 1 GB
- Virtual Address space is 4 GB

Mapping



Page Faults

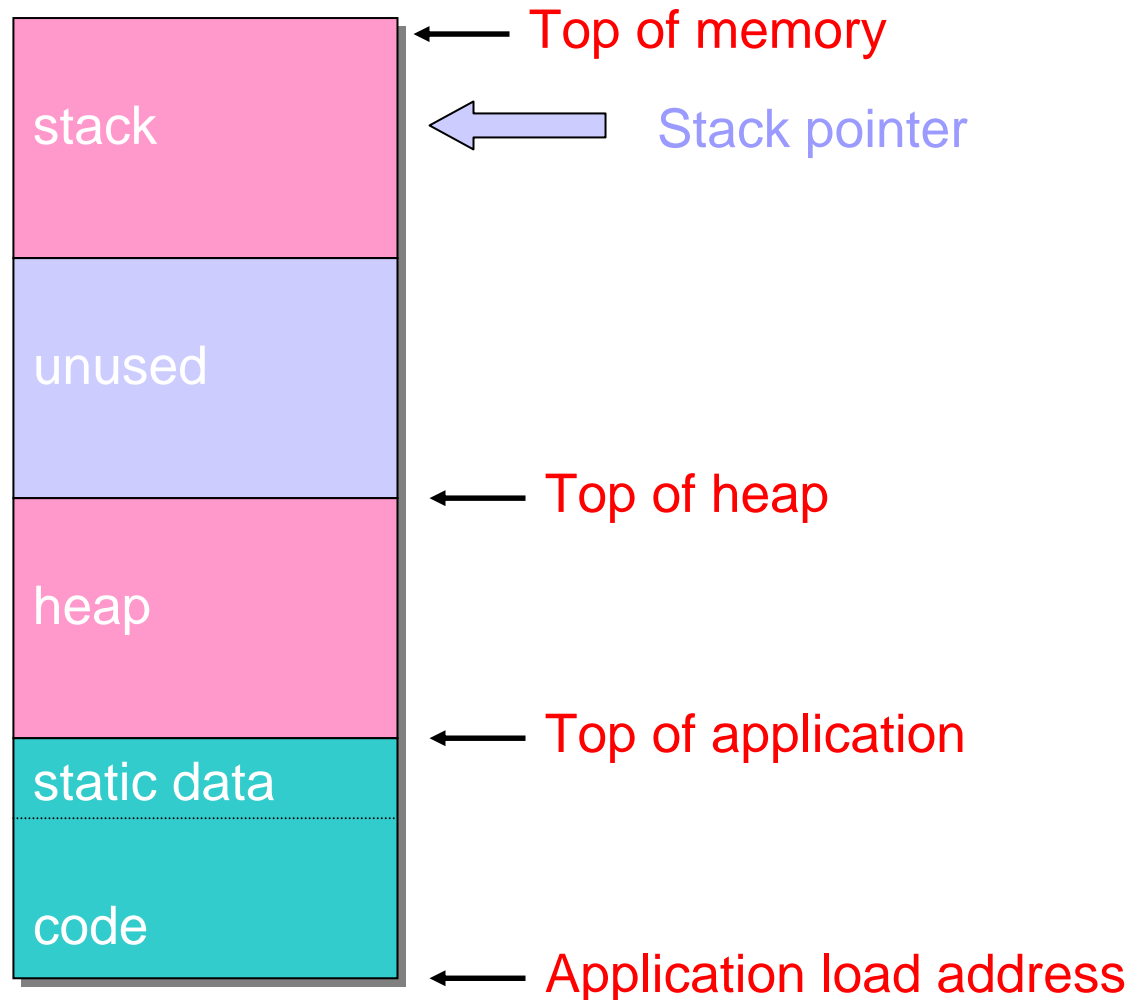
- If the valid bit is off, page fault occurs
- Page faults are very costly, in terms of time
- This results in exception and kernel gets the control
- Kernel finds the page on the disk
- Decides where to put the page in main memory
 - Generally LRU is used
- Updates the page table
- To make translations fast, Translation Look aside Buffer (TLB) is used
- TLB acts like a cache for the page table.

Memory Subsystem of ARM

Memory map

- Arranged as a linear set of logical address
- A 'C' program expects to access
 - A fixed area of program memory
 - A dynamically changing data areas
 - Stack: for local variables and function calls
 - Heap: to satisfy program requests for more memory for data structures
- Application is loaded at lowest address
- Unused memory is allocated on demand to heap or stack

ARM address space model




Data Storage

- Basic data types
 - Char: byte
 - Short int: half word
 - Int, single precision float: word
 - Double precision float: multiple words
- Derived data types like structs, arrays, unions etc. are defined in terms of basic data types
- ARM instruction set is efficient in loading and storing data items when they are properly aligned in memory
- Storing a word to non-word-aligned memory is very inefficient: can take upto 7 ARM instructions

Data Alignment


- ARM C compiler generally aligns data on appropriate boundaries
 - Bytes at byte address
 - Half-words at even byte address
 - Words at four-byte boundaries
- When several data items of different types are declared at the same time, compiler introduces padding

Memory Efficiency



```
Struct S1 (char c;  
int x; short s;  
example1
```

Occupies 3 words



```
Struct S1 (char c;  
short s; int x;  
example2
```

Occupies only 2 words

Memory in ARM based systems

- Wide variety of embedded systems use ARM
 - Memory system requirements vary from
 - Simple memory block with flat address map
 - Using one or more of memory resources
 - Memory resources
 - Multiple types of memories
 - Caches
 - Write buffers
 - Virtual memory and remapping techniques
 - Memory mapped IO

Memory in ARM based systems

- Initialising and controlling memory system
 - Enable cache, configure it for best performance
 - Set up virtual to physical address mapping
 - Restrict access to memory regions
 - Ensure proper access to memory mapped IO
- The standard way to perform memory system control is CP15

Thank You